

Sai Pratheesh

Bengaluru, India | +91 86180 12764 | saipratheesh15@gmail.com
https://www.linkedin.com/in/sai-pratheesh-0b854a343 | https://github.com/saipratheesh07

SUMMARY

Final-year Information Science Engineering student who builds and ships production-shaped AI systems end to end: ingestion, retrieval, evaluation, API, CI and containerised deployment. Comfortable against real customer data and messy third-party sources, and biased toward measuring a system before improving it. Looking for forward deployed, solutions or applied AI engineering work where the job is getting something working in front of a customer.

TECHNICAL SKILLS

Languages	Python, TypeScript, SQL, C++
Backend and infra	FastAPI, Docker, GitHub Actions CI, MongoDB, SQLAlchemy, SQLite
AI and data	Claude API, LangChain, MCP, ChromaDB, retrieval pipelines, LLM-as-judge evaluation
Frontend	React, TypeScript
Practices	Idempotent ingestion, schema design, eval harness design, cost-aware model routing

PROJECTS

RAG Notes Q&A

Python, FastAPI, ChromaDB, Docker, GitHub Actions | github.com/saipratheesh07/rag-notes-qa

- Retrieval-augmented question answering service that answers strictly from an indexed document corpus and returns nothing when the question shares no vocabulary with the corpus, rather than surfacing arbitrary chunks.
- Built the evaluation harness before optimising the retriever. Retrieval precision is 14/14 across the 14 answerable questions, but the set was authored alongside the system and top-5 over a small corpus is a lenient bar, so it is treated as a regression smoke test rather than a benchmark. Next step is expanding to roughly 50 questions written before inspecting retrieval output.
- Found and fixed a substring bug in the LLM-as-judge scorer where the verdict check for FAITHFUL also matched UNFAITHFUL, meaning every answer had been scoring as faithful and the negative test could never fail.
- Retrieval scoring, ingestion and the full test suite run with no API key, so anyone cloning the repository can reproduce the numbers. 13 automated tests green on every push via CI.
- Shipped as a single command Docker deployment, with a .dockerignore added after noticing the image build would otherwise have copied the local .env and its API key into the image.

ChurnLens

Python, SQLAlchemy, Claude API, Apple RSS | *In progress*

- Pipeline that turns public customer reviews into a ranked churn report, comparing a target company against competitors on the same complaint themes.
- Three layer schema separating immutable raw payloads, normalised reviews and model output keyed by model version, so a prompt change re-runs classification only and never re-fetches. Ingestion is idempotent at the database level via a uniqueness constraint on source, entity and content hash.
- Cost-aware classification cascade sends every review to a cheap model first and escalates only low-confidence cases, with escalation rate and cost per thousand reviews reported.
- Running it against a live App Store corpus surfaced two real defects: a source adapter returning zero rows silently, and a theme taxonomy that was hardcoded despite being documented as configurable. Both fixed.

Befit Multi-Agent SEO Content Pipeline

Python, Claude API, WordPress | *Client project*

- Multi-agent pipeline handling keyword research, drafting, editing and publishing to WordPress for a D2C fitness client, with brand voice rules enforced at the editing stage.
- Delivered as a repeatable content system rather than one-off output, so the client could run it without ongoing involvement.

EXPERIENCE

Freelance Engineer, Innova Marketing Group

Bengaluru | 2026

- Building an internal stock and inventory tool against the client's real operational workflow, scoped and delivered directly with the client rather than through a spec.

EDUCATION

B.E. Information Science and Engineering, Nitte Meenakshi Institute of Technology, Bengaluru

Expected 2027

LANGUAGES

English, Kannada, Hindi, Tamil