

Sai Pratheesh

Bengaluru, India | +91 86180 12764 | saipratheesh15@gmail.com
<https://www.linkedin.com/in/sai-pratheesh-0b854a343> | <https://github.com/saipratheesh07>

SUMMARY

Final-year Information Science Engineering student who builds and ships production-shaped AI systems end to end: ingestion, retrieval, evaluation, API, CI and containerised deployment. Comfortable working directly against real customer data and messy third-party sources, and biased toward measuring a system before improving it. Looking for forward deployed, solutions or applied AI engineering work where the job is getting something working in front of a customer.

TECHNICAL SKILLS

Languages	Python, TypeScript, SQL, C++
Backend and infra	FastAPI, Docker, GitHub Actions CI, MongoDB, SQLAlchemy, SQLite
AI and data	Claude API, LangChain, MCP, ChromaDB, retrieval pipelines, LLM-as-judge evaluation
Frontend	React, TypeScript
Practices	Idempotent ingestion, schema design, eval harness design, cost-aware model routing

PROJECTS

RAG Notes Q&A

Python, FastAPI, ChromaDB, Docker, GitHub Actions | github.com/saipratheesh07/rag-notes-qa

- Retrieval-augmented question answering service that answers strictly from an indexed document corpus and refuses when the corpus does not cover the question, verified by a negative test in the eval set.
- Built an evaluation harness before optimising the retriever: 86 percent retrieval precision across 14 answerable questions, plus LLM-as-judge faithfulness scoring on generated answers.
- Retrieval scoring, ingestion and the full test suite run with no API key, so the repository is reproducible by anyone who clones it. 8 automated tests green on every push via CI.
- Shipped as a single command Docker deployment with the index built at image build time.

ChurnLens

Python, SQLAlchemy, Claude API, Apple RSS | *In progress*

- Pipeline that turns public customer reviews into a ranked churn report, comparing a target company against competitors on the same complaint themes.
- Three layer schema separating immutable raw payloads, normalised reviews and model output keyed by model version, so a prompt change re-runs classification only and never re-fetches.
- Ingestion is idempotent at the database level through a uniqueness constraint on source, entity and content hash, making scheduled collection safe to run unsupervised.
- Cost-aware classification cascade routes every review to a cheap model first and escalates only low-confidence cases to a stronger one, with escalation rate and cost per thousand reviews reported.
- Running it against a live App Store corpus surfaced two real defects: a source adapter that returned zero rows silently, and a theme taxonomy that was hardcoded despite being documented as configurable. Both fixed.

Befit Multi-Agent SEO Content Pipeline

Python, Claude API, WordPress | *Client project*

- Multi-agent pipeline handling keyword research, drafting, editing and publishing to WordPress for a D2C fitness client, with brand voice rules enforced at the editing stage.
- Delivered as a repeatable content system rather than one-off output, so the client could run it without ongoing involvement.

EXPERIENCE

Freelance Engineer, Innova Marketing Group

Bengaluru | *2026*

- Building an internal stock and inventory tool against the client's real operational workflow, scoped and delivered directly with the client rather than through a spec.

EDUCATION

B.E. Information Science and Engineering, Nitte Meenakshi Institute of Technology, Bengaluru

Expected 2027

LANGUAGES

English, Tamil, Hindi, Kannada